

# CS 224n Spring 2024: Assignment #3

Eleanor Zhou

**Due Date: April 30th, Tuesday, 4:30 PM PST.**

This assignment is split into two sections: *Neural Machine Translation with RNNs* and *Analyzing NMT Systems*. The first is primarily coding and implementation focused, whereas the second entirely consists of written, analysis questions. If you get stuck on the first section, you can always work on the second as the two sections are independent of each other. Note that the NMT system is more complicated than the neural networks we have previously constructed within this class and takes about **2 hours to train on a GPU**. Thus, we strongly recommend you get started early with this assignment. Finally, the notation and implementation of the NMT system is a bit tricky, so if you ever get stuck along the way, please come to Office Hours so that the TAs can support you.

## 1. Neural Machine Translation with RNNs (45 points)

In Machine Translation, our goal is to convert a sentence from the *source* language (e.g. Mandarin Chinese) to the *target* language (e.g. English). In this assignment, we will implement a sequence-to-sequence (Seq2Seq) network with attention, to build a Neural Machine Translation (NMT) system. In this section, we describe the **training procedure** for the proposed NMT system, which uses a Bidirectional LSTM Encoder and a Unidirectional LSTM Decoder.




Figure 1: Seq2Seq Model with Multiplicative Attention, shown on the third step of the decoder. Hidden states  $\vec{h}_i^{\text{enc}}$  and cell states  $\vec{c}_i^{\text{enc}}$  are defined on the next page.

## Model description (training procedure)

Given a sentence in the source language, we look up the character or word embeddings from an **embeddings matrix**, yielding  $\mathbf{x}_1, \dots, \mathbf{x}_m$  ( $\mathbf{x}_i \in \mathbb{R}^{e \times 1}$ ), where  $m$  is the length of the source sentence and  $e$  is the embedding size. We then feed the embeddings to a **convolutional layer**<sup>1</sup> while maintaining their shapes. We feed the convolutional layer outputs to the **bidirectional encoder**, yielding hidden states and cell states for both the forwards ( $\rightarrow$ ) and backwards ( $\leftarrow$ ) LSTMs. The forwards and backwards versions are concatenated to give hidden states  $\mathbf{h}_i^{\text{enc}}$  and cell states  $\mathbf{c}_i^{\text{enc}}$ :

$$\mathbf{h}_i^{\text{enc}} = [\overleftarrow{\mathbf{h}_i^{\text{enc}}}, \overrightarrow{\mathbf{h}_i^{\text{enc}}}] \text{ where } \mathbf{h}_i^{\text{enc}} \in \mathbb{R}^{2h \times 1}, \overleftarrow{\mathbf{h}_i^{\text{enc}}}, \overrightarrow{\mathbf{h}_i^{\text{enc}}} \in \mathbb{R}^{h \times 1} \quad 1 \leq i \leq m \quad (1)$$

$$\mathbf{c}_i^{\text{enc}} = [\overleftarrow{\mathbf{c}_i^{\text{enc}}}, \overrightarrow{\mathbf{c}_i^{\text{enc}}}] \text{ where } \mathbf{c}_i^{\text{enc}} \in \mathbb{R}^{2h \times 1}, \overleftarrow{\mathbf{c}_i^{\text{enc}}}, \overrightarrow{\mathbf{c}_i^{\text{enc}}} \in \mathbb{R}^{h \times 1} \quad 1 \leq i \leq m \quad (2)$$

We then initialize the **decoder's** first hidden state  $\mathbf{h}_0^{\text{dec}}$  and cell state  $\mathbf{c}_0^{\text{dec}}$  with a linear projection of the encoder's final hidden state and final cell state.<sup>2</sup>

$$\mathbf{h}_0^{\text{dec}} = \mathbf{W}_h [\overleftarrow{\mathbf{h}_1^{\text{enc}}}, \overrightarrow{\mathbf{h}_m^{\text{enc}}}] \text{ where } \mathbf{h}_0^{\text{dec}} \in \mathbb{R}^{h \times 1}, \mathbf{W}_h \in \mathbb{R}^{h \times 2h} \quad (3)$$

$$\mathbf{c}_0^{\text{dec}} = \mathbf{W}_c [\overleftarrow{\mathbf{c}_1^{\text{enc}}}, \overrightarrow{\mathbf{c}_m^{\text{enc}}}] \text{ where } \mathbf{c}_0^{\text{dec}} \in \mathbb{R}^{h \times 1}, \mathbf{W}_c \in \mathbb{R}^{h \times 2h} \quad (4)$$

With the decoder initialized, we must now feed it a target sentence. On the  $t^{\text{th}}$  step, we look up the embedding for the  $t^{\text{th}}$  subword,  $\mathbf{y}_t \in \mathbb{R}^{e \times 1}$ . We then concatenate  $\mathbf{y}_t$  with the *combined-output vector*  $\mathbf{o}_{t-1} \in \mathbb{R}^{h \times 1}$  from the previous timestep (we will explain what this is later down this page!) to produce  $\overline{\mathbf{y}}_t \in \mathbb{R}^{(e+h) \times 1}$ . Note that for the first target subword (i.e. the start token)  $\mathbf{o}_0$  is a zero-vector. We then feed  $\overline{\mathbf{y}}_t$  as input to the decoder.

$$\mathbf{h}_t^{\text{dec}}, \mathbf{c}_t^{\text{dec}} = \text{Decoder}(\overline{\mathbf{y}}_t, \mathbf{h}_{t-1}^{\text{dec}}, \mathbf{c}_{t-1}^{\text{dec}}) \text{ where } \mathbf{h}_t^{\text{dec}} \in \mathbb{R}^{h \times 1}, \mathbf{c}_t^{\text{dec}} \in \mathbb{R}^{h \times 1} \quad (5)$$

$$(6)$$

We then use  $\mathbf{h}_t^{\text{dec}}$  to compute multiplicative attention over  $\mathbf{h}_1^{\text{enc}}, \dots, \mathbf{h}_m^{\text{enc}}$ :

$$\mathbf{e}_{t,i} = (\mathbf{h}_t^{\text{dec}})^T \mathbf{W}_{\text{attProj}} \mathbf{h}_i^{\text{enc}} \text{ where } \mathbf{e}_t \in \mathbb{R}^{m \times 1}, \mathbf{W}_{\text{attProj}} \in \mathbb{R}^{h \times 2h} \quad 1 \leq i \leq m \quad (7)$$

$$\alpha_t = \text{softmax}(\mathbf{e}_t) \text{ where } \alpha_t \in \mathbb{R}^{m \times 1} \quad (8)$$

$$\mathbf{a}_t = \sum_{i=1}^m \alpha_{t,i} \mathbf{h}_i^{\text{enc}} \text{ where } \mathbf{a}_t \in \mathbb{R}^{2h \times 1} \quad (9)$$

$\mathbf{e}_{t,i}$  is a scalar, the  $i$ th element of  $\mathbf{e}_t \in \mathbb{R}^{m \times 1}$ , computed using the hidden state of the decoder at the  $t$ th step,  $\mathbf{h}_t^{\text{dec}} \in \mathbb{R}^{h \times 1}$ , the attention projection  $\mathbf{W}_{\text{attProj}} \in \mathbb{R}^{h \times 2h}$ , and the hidden state of the encoder at the  $i$ th step,  $\mathbf{h}_i^{\text{enc}} \in \mathbb{R}^{2h \times 1}$ .

We now concatenate the attention output  $\mathbf{a}_t$  with the decoder hidden state  $\mathbf{h}_t^{\text{dec}}$  and pass this through a linear layer, tanh, and dropout to attain the *combined-output vector*  $\mathbf{o}_t$ .

<sup>1</sup>Checkout <https://cs231n.github.io/convolutional-networks> for an in-depth description for convolutional layers if you are not familiar

<sup>2</sup>If it's not obvious, think about why we regard  $[\overleftarrow{\mathbf{h}_1^{\text{enc}}}, \overrightarrow{\mathbf{h}_m^{\text{enc}}}]$  as the 'final hidden state' of the Encoder.

$$\mathbf{u}_t = [\mathbf{a}_t; \mathbf{h}_t^{\text{dec}}] \text{ where } \mathbf{u}_t \in \mathbb{R}^{3h \times 1} \quad (10)$$

$$\mathbf{v}_t = \mathbf{W}_u \mathbf{u}_t \text{ where } \mathbf{v}_t \in \mathbb{R}^{h \times 1}, \mathbf{W}_u \in \mathbb{R}^{h \times 3h} \quad (11)$$

$$\mathbf{o}_t = \text{dropout}(\tanh(\mathbf{v}_t)) \text{ where } \mathbf{o}_t \in \mathbb{R}^{h \times 1} \quad (12)$$

Then, we produce a probability distribution  $\mathbf{P}_t$  over target subwords at the  $t^{\text{th}}$  timestep:

$$\mathbf{P}_t = \text{softmax}(\mathbf{W}_{\text{vocab}} \mathbf{o}_t) \text{ where } \mathbf{P}_t \in \mathbb{R}^{V_t \times 1}, \mathbf{W}_{\text{vocab}} \in \mathbb{R}^{V_t \times h} \quad (13)$$

Here,  $V_t$  is the size of the target vocabulary. Finally, to train the network we then compute the cross entropy loss between  $\mathbf{P}_t$  and  $\mathbf{g}_t$ , where  $\mathbf{g}_t$  is the one-hot vector of the target subword at timestep  $t$ :

$$J_t(\theta) = \text{CrossEntropy}(\mathbf{P}_t, \mathbf{g}_t) \quad (14)$$

Here,  $\theta$  represents all the parameters of the model and  $J_t(\theta)$  is the loss on step  $t$  of the decoder. Now that we have described the model, let's try implementing it for Mandarin Chinese to English translation!

## Setting up Your Local Development Environment

Ensure that you have [conda](#) installed. Unzip the starter code, and open the resulting directory in your favorite python integrated development environment ([Visual Studio Code](#) is a popular choice). Open a terminal and navigate (cd) to the directory that contains the env-cpu.yml and env-gpu.yml files. Create and activate the cs224n-cpu conda environment as follows:

```
conda env create --file env-cpu.yml
conda activate cs224n-cpu
```

## Setting up Your Cloud GPU-powered Virtual Machine

Follow the instructions in the [GCP Guide for CS224n](#) (link also provided on website and Ed) to create your VM instance. This should take you approximately 45 minutes. Though you will need the GPU to train your model, we strongly advise that you first develop the code locally and ensure that it runs, before attempting to train it on your VM. GPU time is expensive and limited. It takes approximately **1.5 to 2 hours** to train the NMT system. We don't want you to accidentally use all your GPU time for debugging your model rather than training and evaluating it. Finally, **make sure that your VM is turned off whenever you are not using it.**

If your GCP subscription runs out of money, your VM will be temporarily locked and inaccessible. If that happens, please make a private Ed post describing the situation.

After setting up the cloud VM, you should turn it off while you work on the programming assignment locally.

## Implementation and written questions

- (a) (2 points) (coding) In order to apply tensor operations, we must ensure that the sentences in a given batch are of the same length. Thus, we must identify the longest sentence in a batch and pad others to be the same length. Implement the `pad_sents` function in `utils.py`, which shall produce these padded sentences.
- (b) (3 points) (coding) Implement the `__init__` function in `model_embeddings.py` to initialize the necessary source and target embeddings.
- (c) (4 points) (coding) Implement the `__init__` function in `nmt_model.py` to initialize the necessary model layers (LSTM, CNN, projection, and dropout) for the NMT system.
- (d) (8 points) (coding) Implement the `encode` function in `nmt_model.py`. This function converts the padded source sentences into the tensor  $\mathbf{X}$ , generates  $\mathbf{h}_1^{\text{enc}}, \dots, \mathbf{h}_m^{\text{enc}}$ , and computes the initial state  $\mathbf{h}_0^{\text{dec}}$  and initial cell  $\mathbf{c}_0^{\text{dec}}$  for the Decoder. You can run a non-comprehensive sanity check by executing:

```
python sanity_check.py 1d
```

- (e) (8 points) (coding) Implement the `decode` function in `nmt_model.py`. This function constructs  $\bar{\mathbf{y}}$  and runs the step function over every timestep for the input. You can run a non-comprehensive sanity check by executing:

```
python sanity_check.py 1e
```

- (f) (10 points) (coding) Implement the `step` function in `nmt_model.py`. This function applies the Decoder's LSTM cell for a single timestep, computing the encoding of the target subword  $\mathbf{h}_t^{\text{dec}}$ , the attention scores  $\mathbf{e}_t$ , attention distribution  $\alpha_t$ , the attention output  $\mathbf{a}_t$ , and finally the combined output  $\mathbf{o}_t$ . You can run a non-comprehensive sanity check by executing:

```
python sanity_check.py 1f
```

- (g) (3 points) (written) The `generate_sent_masks()` function in `nmt_model.py` produces a tensor called `enc_masks`. It has shape (batch size, max source sentence length) and contains 1s in positions corresponding to 'pad' tokens in the input, and 0s for non-pad tokens. Look at how the masks are used during the attention computation in the `step()` function (lines 295-296).

First explain (in around three sentences) what effect the masks have on the entire attention computation. Then explain (in one or two sentences) why it is necessary to use the masks in this way.

Now it's time to get things running! As noted earlier, we recommend that you develop the code on your personal computer. Confirm that you are running in the proper conda environment and then execute the following command to train the model on your local machine:

```
sh run.sh train_local
(Windows) run.bat train_local
```

To help with monitoring and debugging, the starter code uses tensorboard to log loss and perplexity during training using TensorBoard<sup>3</sup>. TensorBoard provides tools for logging and visualizing training information from experiments. To open TensorBoard and monitor the training process, run the following in a separate terminal in which you have also activated the `cs224n-cpu` conda environment, then access tensorboard at <http://localhost:6006/>.

```
tensorboard --logdir=runs --port 6006
```




Figure 2: Tensorboard showing loss and perplexity values on your local machine

You should see a significant decrease in loss during the initial iterations (Fig 2). Once your code runs for a few hundreds of iterations without crashing, power on your VM from the GCP Console.

**Train the NMT system on the VM:** Refer to the GCP How-to Appendix to:

- Setup your GCP VM and Obtain SSH Connection Info
  - Connect to the VM (with SSH Tunneling setup so that you can view remote tensorboard logs)
  - Copy your code from your computer to the cloud VM
  - Train the NMT System on the Cloud VM
  - Download the Gradescope Submission Package from Your Cloud VM
- (h) (3 points) (written) Once your model is done training (**this should take under 2 hours on the VM**), execute the following command to test the model:

```
sh run.sh test
```

Please report the model's corpus BLEU Score. It should be larger than 18.

- (i) (4 points) (written) In class, we learned about dot product attention, multiplicative attention, and additive attention. As a reminder, dot product attention is  $\mathbf{e}_{t,i} = \mathbf{s}_t^T \mathbf{h}_i$ , multiplicative attention is  $\mathbf{e}_{t,i} = \mathbf{s}_t^T \mathbf{W} \mathbf{h}_i$ , and additive attention is  $\mathbf{e}_{t,i} = \mathbf{v}^T \tanh(\mathbf{W}_1 \mathbf{h}_i + \mathbf{W}_2 \mathbf{s}_t)$ .
- (2 points) Explain one advantage and one disadvantage of *dot product attention* compared to multiplicative attention.
  - (2 points) Explain one advantage and one disadvantage of *additive attention* compared to multiplicative attention.

<sup>3</sup><https://pytorch.org/docs/stable/tensorboard.html>

(g) The sentence mask sets the attention scores  $e_{t,i}$  at padded source positions to  $-\infty$  before the softmax is applied. Consequently, the corresponding attention weights  $\alpha_{t,i}$  become zero, while the weights over the non-padding positions are renormalized to sum to one. Therefore, the attention output  $\mathbf{a}_t$  is computed using only the encoder hidden states corresponding to actual source tokens, and the subsequent combined output and target-word distribution are not influenced by padded positions. This masking is necessary because padding tokens are artifacts introduced only to construct equal-length batches and contain no linguistic information. Without masking, the decoder could assign attention mass to meaningless padded positions, thereby distorting the source representation.

(h) The model achieves a corpus BLEU score of 21.38.

(i)

- (i) One advantage of dot-product attention is that it does not introduce an additional learned projection matrix, so it uses fewer parameters and is computationally cheaper than multiplicative attention. A disadvantage is that the encoder and decoder hidden states must have the same dimensionality. Moreover, dot-product attention cannot learn a transformation that models a more flexible compatibility relation between the encoder and decoder representations.
- (ii) One advantage of additive attention is that its learned nonlinear compatibility function is more expressive and can naturally compare encoder and decoder states with different dimensionalities. A disadvantage is that it requires additional projection matrices, a parameter vector, and a nonlinear activation for every source–target pair. It is therefore generally more computationally expensive than multiplicative attention.

## 2. Analyzing NMT Systems (25 points)

- (a) (3 points) Look at the `src.vocab` file for some examples of phrases and words in the source language vocabulary. When encoding an input Mandarin Chinese sequence into “pieces” in the vocabulary, the tokenizer maps the sequence to a series of vocabulary items, each consisting of one or more characters (thanks to the `sentencepiece` tokenizer, we can perform this segmentation even when the original text has no white space). Given this information, how could adding a 1D Convolutional layer after the embedding layer and before passing the embeddings into the bidirectional encoder help our NMT system? **Hint:** each Mandarin Chinese character is either an entire word or a morpheme in a word. Look up the meanings of 电, 脑, and 电脑 separately for an example. The characters 电 (electricity) and 脑 (brain) when combined into the phrase 电脑 mean computer.
- (b) (8 points) Here we present a series of errors we found in the outputs of our NMT model (which is the same as the one you just trained). For each example of a reference (i.e., ‘gold’) English translation, and NMT (i.e., ‘model’) English translation, please:
1. Identify the error in the NMT translation.
  2. Provide possible reason(s) why the model may have made the error (either due to a specific linguistic construct or a specific model limitation).
  3. Describe one possible way we might alter the NMT system to fix the observed error. There are more than one possible fixes for an error. For example, it could be tweaking the size of the hidden layers or changing the attention mechanism.

Below are the translations that you should analyze as described above. Only analyze the underlined error in each sentence. Rest assured that you don’t need to know Mandarin to answer these questions. You just need to know English! If, however, you would like some additional color on the source sentences, feel free to use a resource like [https://www.archchinese.com/chinese\\_english\\_dictionary.html](https://www.archchinese.com/chinese_english_dictionary.html) to look up words. Feel free to search the training data file to have a better sense of how often certain characters occur.

- (2 points) **Source Sentence:** 贼人其后被警方拘捕及被判处盗窃罪名成立。  
**Reference Translation:** the culprits were subsequently arrested and convicted.  
**NMT Translation:** the culprit was subsequently arrested and sentenced to theft.
- (2 points) **Source Sentence:** 几乎已经没有地方容纳这些人, 资源已经用尽。  
**Reference Translation:** there is almost no space to accommodate these people, and resources have run out.  
**NMT Translation:** the resources have been exhausted and resources have been exhausted.
- (2 points) **Source Sentence:** 当局已经宣布今天是国殇日。  
**Reference Translation:** authorities have announced a national mourning today.  
**NMT Translation:** the administration has announced today's day.
- (2 points) **Source Sentence**<sup>4</sup>: 俗语有云:“唔做唔错”。  
**Reference Translation:** “act not, err not”, so a saying goes.  
**NMT Translation:** as the saying goes, “it’s not wrong.”

- (c) (14 points) BLEU score is the most commonly used automatic evaluation metric for NMT systems. It is usually calculated across the entire test set, but here we will consider BLEU defined for a single example.<sup>5</sup> Suppose we have a source sentence  $s$ , a set of  $k$  reference translations  $\mathbf{r}_1, \dots, \mathbf{r}_k$ , and a

<sup>4</sup>This is a Cantonese sentence! The data used in this assignment comes from GALE Phase 3, which is a compilation of news written in simplified Chinese from various sources scraped from the internet along with their translations. For more details, see <https://catalog.ldc.upenn.edu/LDC2017T02>.

<sup>5</sup>This definition of sentence-level BLEU score matches the `sentence_bleu()` function in the `nlTK` Python package. Note that the `NLTK` function is sensitive to capitalization. In this question, all text is lowercased, so capitalization is irrelevant. [http://www.nltk.org/api/nltk.translate.html#nltk.translate.bleu\\_score.sentence\\_bleu](http://www.nltk.org/api/nltk.translate.html#nltk.translate.bleu_score.sentence_bleu)

candidate translation  $\mathbf{c}$ . To compute the BLEU score of  $\mathbf{c}$ , we first compute the *modified  $n$ -gram precision*  $p_n$  of  $\mathbf{c}$ , for each of  $n = 1, 2, 3, 4$ , where  $n$  is the  $n$  in **n-gram**:

$$p_n = \frac{\sum_{\text{ngram} \in \mathbf{c}} \min \left( \max_{i=1, \dots, k} \text{Count}_{\mathbf{r}_i}(\text{ngram}), \text{Count}_{\mathbf{c}}(\text{ngram}) \right)}{\sum_{\text{ngram} \in \mathbf{c}} \text{Count}_{\mathbf{c}}(\text{ngram})} \quad (15)$$

Here, for each of the  $n$ -grams that appear in the candidate translation  $\mathbf{c}$ , we count the maximum number of times it appears in any one reference translation, capped by the number of times it appears in  $\mathbf{c}$  (this is the numerator). We divide this by the number of  $n$ -grams in  $\mathbf{c}$  (denominator).

Next, we compute the *brevity penalty* BP. Let  $\text{len}(\mathbf{c})$  be the length of  $\mathbf{c}$  and let  $\text{len}(\mathbf{r})$  be the length of the reference translation that is closest to  $\text{len}(\mathbf{c})$  (in the case of two equally-close reference translation lengths, choose  $\text{len}(\mathbf{r})$  as the shorter one).

$$BP = \begin{cases} 1 & \text{if } \text{len}(\mathbf{c}) \geq \text{len}(\mathbf{r}) \\ \exp \left( 1 - \frac{\text{len}(\mathbf{r})}{\text{len}(\mathbf{c})} \right) & \text{otherwise} \end{cases} \quad (16)$$

Lastly, the BLEU score for candidate  $\mathbf{c}$  with respect to  $\mathbf{r}_1, \dots, \mathbf{r}_k$  is:

$$BLEU = BP \times \exp \left( \sum_{n=1}^4 \lambda_n \log p_n \right) \quad (17)$$

where  $\lambda_1, \lambda_2, \lambda_3, \lambda_4$  are weights that sum to 1. The log here is natural log.

- i. (5 points) Please consider this example:

Source Sentence  $\mathbf{s}$ : 需要有充足和可预测的资源。

Reference Translation  $\mathbf{r}_1$ : *resources have to be sufficient and they have to be predictable*

Reference Translation  $\mathbf{r}_2$ : *adequate and predictable resources are required*

NMT Translation  $\mathbf{c}_1$ : there is a need for adequate and predictable resources

NMT Translation  $\mathbf{c}_2$ : resources be sufficient and predictable to

Please compute the BLEU scores for  $\mathbf{c}_1$  and  $\mathbf{c}_2$ . Let  $\lambda_i = 0.5$  for  $i \in \{1, 2\}$  and  $\lambda_i = 0$  for  $i \in \{3, 4\}$  (**this means we ignore 3-grams and 4-grams**, i.e., don't compute  $p_3$  or  $p_4$ ).

When computing BLEU scores, show your work (i.e., show your computed values for  $p_1$ ,  $p_2$ ,  $\text{len}(\mathbf{c})$ ,  $\text{len}(\mathbf{r})$  and  $BP$ ). Note that the BLEU scores can be expressed between 0 and 1 or between 0 and 100. The code is using the 0 to 100 scale while in this question we are using the **0 to 1** scale. Please round your responses to 3 decimal places.

Which of the two NMT translations is considered the better translation according to the BLEU Score? Do you agree that it is the better translation?

- ii. (5 points) Our hard drive was corrupted and we lost Reference Translation  $\mathbf{r}_1$ . Please recompute BLEU scores for  $\mathbf{c}_1$  and  $\mathbf{c}_2$ , this time with respect to  $\mathbf{r}_2$  only. Which of the two NMT translations now receives the higher BLEU score? Do you agree that it is the better translation?
- iii. (2 points) Due to data availability, NMT systems are often evaluated with respect to only a single reference translation. Please explain (in a few sentences) why this may be problematic. In your explanation, discuss how the BLEU score metric assesses the quality of NMT translations when there are multiple reference translations versus a single reference translation.
- iv. (2 points) List two advantages and two disadvantages of BLEU, compared to human evaluation, as an evaluation metric for Machine Translation.




Figure 3: Translation with beam search results for an example sentence are recorded in tensorboard for various iterations. The same data is available in the `outputs/beam_search_diagnostics/` folder in your working directory.

- (d) (4 points) *Beam search* is often employed to improve the quality of machine translation systems. While you were training the model, beam search results for the same example sentence at different iterations were also recorded in TensorBoard, and accessible in the *TEXT* tab (Fig 3).

The recorded diagnostic information includes json documents with the following fields: `example_source` (the source sentence tokens), `example_target` (the ground truth target sentence tokens), and `hypotheses` (10 hypotheses corresponding to the search result with beam size 10). Note that a predicted translation is often called *hypothesis* in the neural machine translation jargon.

- i. (2 points) Did the translation quality improve over the training iterations for the model? Give three examples of translations of the example sentence at iterations 200, 3000, and the last iteration to illustrate your answer. For each iteration, pick the first beam search hypothesis as an example:
- ii. (2 points) How do various hypotheses resulting from beam search qualitatively compare? Give three other examples of hypotheses proposed by beam search at the last iteration to illustrate your answer.

(a) A 1d convolutional layer can combine the embeddings of adjacent SentencePiece tokens and extract local character or morpheme patterns before they are passed to the bidirectional encoder. This is particularly useful in Chinese because a multi-character word may have a meaning that cannot be obtained by independently interpreting its individual characters. For example, 电 and 脑 separately mean ‘electricity’ and ‘brain,’ while the combination 电脑 means ‘computer.’ A convolution with a kernel spanning neighboring tokens can learn a representation for such local combinations, thereby reducing the burden on the recurrent encoder to reconstruct word-level meanings from independently embedded pieces.

(b)

- (i) **Error:** The model translates the plural phrase *the culprits were* as the singular phrase *the culprit was*.

**Possible reason:** Mandarin nouns generally do not explicitly mark grammatical number, so the model may not have enough information to decide whether the noun should be singular or plural in English. The decoder may therefore favor the more frequent singular construction.

**Possible fix:** The model could be trained with additional number-sensitive examples or data augmentation in which singular and plural English constructions are explicitly contrasted. Linguistic number features could also be provided to the encoder or decoder.

- (ii) **Error:** The model repeats *resources have been exhausted* and completely omits the clause stating that there is almost no space to accommodate the people.

**Possible reason:** This is an under-translation and repetition error. The attention mechanism may repeatedly attend to the source region corresponding to ‘resources have been exhausted’ while failing to cover the source tokens corresponding to the first clause.

**Possible fix:** A coverage-based attention mechanism or a coverage penalty during beam search could keep track of which source positions have already been translated. This would discourage repeated attention to the same phrase and encourage the decoder to cover the omitted source clause.

- (iii) **Error:** The phrase *national mourning day* is incorrectly translated as the semantically meaningless phrase *today’s day*.

**Possible reason:** The source contains a relatively uncommon multi-character expression whose meaning is not well represented by the meanings of its individual subword pieces. Limited occurrences of this phrase in the training data may prevent the model from learning its idiomatic meaning.

**Possible fix:** The model could be trained or fine-tuned on more parallel examples containing this expression and related compounds. Pretrained contextual representations or an external phrase-level lexicon could also improve its treatment of rare multi-character expressions.

- (iv) **Error:** The idiom meaning approximately ‘if one does nothing, one makes no mistakes’ is reduced to *it’s not wrong*, which loses both the parallel negation structure and the intended meaning.

**Possible reason:** The source sentence is Cantonese, whereas most of the training data is written in standard Mandarin. The Cantonese tokens and the idiomatic construction are therefore likely rare or out-of-distribution for the model.

**Possible fix:** The model could be fine-tuned on Cantonese–English parallel data, especially data containing Cantonese idioms. A dialect-aware multilingual model or an external idiom lexicon could also help preserve the intended meaning.

(c)

- (i) For  $\mathbf{c}_1$ , the candidate length is  $\text{len}(\mathbf{c}_1) = 9$ . The reference lengths are  $\text{len}(\mathbf{r}_1) = 11$  and  $\text{len}(\mathbf{r}_2) = 6$ , so the closest reference is  $\mathbf{r}_1$  and

$$BP_1 = \exp\left(1 - \frac{11}{9}\right) \approx 0.801.$$

Four of the nine candidate unigrams occur in at least one reference, and three of its eight bigrams occur in at least one reference. Therefore,

$$p_1(\mathbf{c}_1) = \frac{4}{9} \approx 0.444, \quad p_2(\mathbf{c}_1) = \frac{3}{8} = 0.375.$$

Its BLEU score is

$$\begin{aligned} BLEU(\mathbf{c}_1) &= 0.801 \exp\left(\frac{1}{2} \log \frac{4}{9} + \frac{1}{2} \log \frac{3}{8}\right) \\ &\approx 0.327. \end{aligned}$$

For  $\mathbf{c}_2$ ,  $\text{len}(\mathbf{c}_2) = 6$ , which is equal to the length of  $\mathbf{r}_2$ , so  $BP_2 = 1$ . All six candidate unigrams occur in at least one reference, and three of its five bigrams occur in at least one reference. Hence,

$$p_1(\mathbf{c}_2) = \frac{6}{6} = 1, \quad p_2(\mathbf{c}_2) = \frac{3}{5} = 0.600,$$

and

$$BLEU(\mathbf{c}_2) = \exp\left(\frac{1}{2} \log 1 + \frac{1}{2} \log \frac{3}{5}\right) \approx 0.775.$$

Thus, BLEU considers  $\mathbf{c}_2$  to be the better translation. I do not agree with this ranking. Although  $\mathbf{c}_2$  has high lexical overlap with the references, it is ungrammatical and does not clearly express that sufficient and predictable resources are required. In contrast,  $\mathbf{c}_1$  is fluent and conveys the intended meaning more accurately.

- (ii) When only  $\mathbf{r}_2$  is available,  $\mathbf{c}_1$  retains

$$p_1(\mathbf{c}_1) = \frac{4}{9}, \quad p_2(\mathbf{c}_1) = \frac{3}{8}.$$

Since  $\text{len}(\mathbf{c}_1) = 9 > \text{len}(\mathbf{r}_2) = 6$ , its brevity penalty is now 1. Therefore,

$$BLEU(\mathbf{c}_1) = \exp\left(\frac{1}{2} \log \frac{4}{9} + \frac{1}{2} \log \frac{3}{8}\right) \approx 0.408.$$

For  $\mathbf{c}_2$ , only three of its six unigrams and one of its five bigrams occur in  $\mathbf{r}_2$ . Its length is equal to the reference length, so its brevity penalty remains 1. Thus,

$$p_1(\mathbf{c}_2) = \frac{3}{6} = 0.500, \quad p_2(\mathbf{c}_2) = \frac{1}{5} = 0.200,$$

and

$$BLEU(\mathbf{c}_2) = \exp\left(\frac{1}{2} \log \frac{1}{2} + \frac{1}{2} \log \frac{1}{5}\right) \approx 0.316.$$

With only  $\mathbf{r}_2$ ,  $\mathbf{c}_1$  receives the higher BLEU score. I agree with this ranking because  $\mathbf{c}_1$  is fluent and semantically expresses the requirement for adequate and predictable resources, whereas  $\mathbf{c}_2$  is ungrammatical and incomplete.

- (iii) A source sentence can have many correct translations that use different words, syntactic structures, or paraphrases. With only one reference, BLEU may assign a low score to a valid translation simply because its surface  $n$ -grams differ from those of that particular reference. Multiple references increase the set of acceptable matching  $n$ -grams, making the score less sensitive to one annotator's wording and reducing the chance that a good translation is unfairly penalized.
- (iv) Two advantages of BLEU are that it is inexpensive and fast to compute, and that it is automatic and reproducible, allowing large systems to be compared consistently. Two disadvantages are that it primarily measures surface  $n$ -gram overlap rather than semantic correctness or fluency, and that its score is highly dependent on the number and wording of the available reference translations. Consequently, BLEU can disagree with human judgments, especially at the sentence level.

## Submission Instructions

You shall submit this assignment on GradeScope as two submissions – one for “Assignment 3 [coding]” and another for “Assignment 3 [written]”:

1. Run the `collect_submission.sh` script on the cloud VM to produce your `assignment3.zip` file. See *How to Download the Gradescope Submission Package from Your Cloud VM* in the GCP How-to Appendix.
2. Upload your `assignment3.zip` file to GradeScope to “Assignment 3 [coding]”.
3. Upload your written solutions to GradeScope to “Assignment 3 [written]”. When you submit your assignment, make sure to tag all the pages for each problem according to Gradescope's submission directions. Points will be deducted if the submission is not correctly tagged.

## Appendix: GCP How-to

**How to Obtain SSH connection info** After following the instructions in the [GCP Guide for CS224n](#), your Google Cloud Compute Engine dashboard should look like Fig 4. Click on the arrow button near SSH and select *View gcloud command* as shown in the figure. What you will get looks like the following. Make a note of the zone, the machine name, and the project name which in the example are respectively us-west4-b, nvidia-gpu-optimized-vmi-1-vm and grand-hangar-420500. **Yours will be different!**

```
gcloud compute ssh --zone "us-west4-b" "nvidia-gpu-optimized-vmi-1-vm" --
  ↪ project "grand-hangar-420500"
```




Figure 4: The cloud VM you create should be based on the [nvidia-gpu-optimized-vmi](#) virtual machine image (VMI). After deploying the instance, click on the menu near SSH and select *View gcloud command* to obtain the SSH connection information.

**How to Connect to the Cloud VM with SSH Tunelling** The following ssh connection command, which includes the `--ssh-flag "-L 6007:localhost:6007"` option, will also forward port 6007 of the cloud machine to the same port number on your local machine (ssh tunneling). Therefore, if the tensorboard server is running on the cloud VM and listening to port 6007, after establishing an ssh connection as follows, you will be able to access at <http://localhost:6007/> in a web browser on your local machine. Run the following ssh command to connect to the cloud machine. **be sure to use your own zone, project name and machine name**

```
gcloud compute ssh --zone "us-west4-b" "nvidia-gpu-optimized-vmi-1-vm" --
  ↪ project "grand-hangar-420500" --ssh-flag "-L 6007:localhost:6007"
```

### How to Copy Your Code From Your Computer to the Cloud VM

- zip current directory, which should contain the env-cpu.yml and env-gpu.yml files.

```
zip -r student.zip *
```

- copy the obtained zip file to the cloud machine

```
gcloud compute scp student.zip nvidia-gpu-optimized-vmi-1-vm:~/ --zone "
  ↪ us-west4-b" --project "grand-hangar-420500"
```

## How to Train the NMT System on the Cloud VM

- If zip is not installed on your cloud machine. Install it as follows and try again.

```
sudo apt-get install zip
```

- Unzip student folder on the cloud machine

```
unzip student.zip -d student
cd student
```

- Create conda GPU environment

```
conda env create --file env-gpu.yml
```

- Activate conda GPU environment

```
conda activate cs224n-nmt-gpu
```

- create a tmux session in which you start training the machine translation model

```
tmux new -s s-nmt
```

```
# start the training (in the tmux session)
sh run.sh train
```

```
# detach from the tmux session
CTRL+B D
```

- create a tmux session in which you start TensorBoard

```
tmux new -s s-tboard
```

```
# start the tensorboard (in the tmux session)
tensorboard --logdir runs/ --port 6007
```

```
# detach from the tmux session
CTRL+B D
```

- You can attach/detach from the two tmux sessions as needed. See the appendix on Tmux for basic use cases.

- If you established the SSH connection with the appropriate port forwarding options, you can view tensorboard in a web browser on your local machine at this address <http://localhost:6007/> (Fig. 5)

- when the training is complete, run the test phase

```
tmux attach -t s-nmt
sh run.sh test
```

- The next section shows how to make a grade scope submission with the artifacts generated on the cloud VM.




Figure 5: Tensorboard running on your local machine and showing loss and perplexity values on your cloud virtual machine.

### How to Download the Gradescope Submission Package from Your Cloud VM

- Connect to the cloud VM as specified above
- On the Cloud VM, run the following command to create the gradescope submission package (assignment3.zip)

```
sh collect_submission.sh
```

- Disconnect from the cloud VM (by running **exit**), or open a terminal on your local machine.
- Download the GradeScope submission package to your local machine

```
gcloud compute scp nvidia-gpu-optimized-vmi-1-vm:~/student/assignment3.zip
↪ . --zone "us-west4-b" --project "grand-hangar-420500"
```

- Submit assignment3.zip to GradeScope

## Appendix: Tmux How-to

tmux (terminal multiplexer) is a tool that lets you create and toggle between multiple terminals that run in the background, and keep them running even after you disconnect your ssh session. The following assumes that you are already connected (via ssh) to your cloud virtual machine instance, and shows you basic use cases of tmux. Learn more about tmux at <https://github.com/tmux/tmux/wiki>

- Create and attach to a new tmux session (named *s-tboard*)

```
tmux new -s s-tboard
```

- Run a program in the new session. We will run tensorboard

```
tensorboard --logdir runs/ --port 6007
```

- (From within the tmux session), Detach from a tmux session, and leave the program you started running in the session.

```
CTRL+b d
```

- (After detaching from tmux), view all tmux sessions

```
tmux ls
```

- attach to a particular tmux session (named s-nmt)

```
tmux attach -t s-nmt
```

- to terminate a tmux session (from within the session)

```
exit
```